A Simple Survey on Diffusion Models for Image Synthesis Since 2020

Zhixin Ling

ABSTRACT

Image Synthesis is a basic research task in computer vision. In the early deep learning era, GANs, VAEs, ARs and flow models are comprehensively adopted for this task. Recently, diffusion models (DMs) are shown to be able to generate various high-quality images. Hence, the relevant research on DM has become a hotspot. In this paper, we present a simple survey on DMs since 2020. We focus on the application of Image Synthesis and closely related works. We will show HOW these diffusion models work but WHY is not our focus. Detailed mathematical derivations will be skipped. Similarities and differences of these works will be emphasized.

1 INTRODUCTION

Our paper mainly consists of three parts:

(1) Horizontal Comparison (Sec.2). We briefly look through deep image synthesis methods and show their advantages and differences.
(2) Vertical Comparison (Sec.3). We present comparisons between diffusion models since 2020. We are going to discuss the differences between their motivations in a high-level view. Also, their open-source repositories and datasets will be summarized.
(3) Conclusion (Sec.4). We conclude our investigation on diffusion models and present some ideas for future research.

2 HORIZONTAL COMPARISON

Neural networks are able to learn various functions, but do NOT directly learn probabilistic distributions. As a result, deep generative models parameterize various complex probabilistic distributions and learn the distribution parameters, *e.g.*, mean and variance of a Gaussian distribution.

Deep image synthesis models can be roughly categorized into Five groups: generative adversarial networks (GANs), autoregressive models (ARs), flows, variational autoencoders (VAEs), diffusion models (DMs). Their main differences are listed as follows:

(1) VAE [7] encodes an input image into latent variables, which follows a Gaussian distribution.

(2) GAN [3] adopts a generator to synthetize an image and a discriminator to distinguish a true image from a false one.
(3) Flow [2] uses a reversible function f to encode an input image into latent variables and decode the latent variables via reversed f.
(4) AR [15] generates an image pixel by pixel, like tangling a next-word prediction task in NLP.
(5) DM [5] learns decodes latent variables into an image step by step. In each step, new latent variables are sampled from a Gaussian distribution that is depended on the current latent variables.

3 VERTICAL COMPARISON

In these part, we introduce different diffusion models. We try to skip heavy mathematical theorems and directly go to research conclusions. Also, for simplicity, we will ignore some unimportant details, *e.g.*, a constant scale parameter.

3.1 DDPM

Denoising Diffusion Probabilistic Model(DDPM)[5] is a pioneer work. To make this work clear is crucial.

3.1.1 Forward Process. DDPM encodes an given image x_0 into Gaussian noises x_T step by step, which is named a **forward** process. According to the chain rule, we explain the process in formula:

$$q(x_t|x_{t-1}) = \mathcal{N}(\sqrt{1-\beta_t}x_{t-1}, \beta_t \mathbf{I});$$
$$q(x_T|x_0) = \prod_{t=1}^{T} q(x_t|x_{t-1}). \tag{1}$$

where β_t ranges from 0 to 1 and grows with t. Therefore, Gaussian noises employed on x_t step by step. At last, $q(x_T|x_{T-1})$ obeys a Gaussian distribution that is not related with x_{T-1} at all.

3.1.2 Backward Process. The decoding stage, also a **backward** process, is performed inversely:

$$p(x_T) = \mathcal{N}(0, \mathbf{I});$$
$$q(x_{t-1}|x_t) = \mathcal{N}(\mu(x_t, t), \Sigma(x_t, t));$$
$$p(x_0) = p(x_T) \prod_{t=T}^{0} q(x_{t-1}|x_t). \tag{2}$$

In this process, μ and Σ are modeled by neural networks. The backward process starts from $\mathcal{N}(0, \mathbf{I})$. In each step, a Gaussian distribution is deriviated according to the current noised image x_t. The the $(t-1)$-step image is sampled from the new distribution. Therefore, the image synthesis process is actually converted into a denoising process, just as indicated by the model title **Denoising** Diffusion Probabilistic Model.

In test stage (*i.e.*, synthesis stage), the backward process is used for image synthesis.

Basic Model	A Simplified Naive Loss	Generation Example
VAE [7] (ICLR'14)	$\mathcal{L}_{KL} = KL(q(\mathbf{z}\|x)\|p(\mathbf{z})), \mathcal{L}_{rec} = \|p(x\|\mathbf{z}) - p(x)\|^2, p(\mathbf{z}) = \mathcal{N}(0,1)$	$\hat{x} = p(x\|\mathbf{z}), \mathbf{z} \sim \mathcal{N}(0,1)$
GAN [3] (NIPS'14)	$\mathcal{L}_D = \log D(x) + \log(1 - D(G(\mathbf{z}))), \mathcal{L}_G = \log(D(G(\mathbf{z}))), \mathbf{z} \sim \mathcal{N}(0,1)$	$\hat{x} = G(\mathbf{z}), \mathbf{z} \sim \mathcal{N}(0,1)$
Flow [2] (ICLR'14)	$\mathcal{L}_{flow} = \|x - f(\mathbf{z})\|^2 + \|\mathbf{z} - f^{-1}(x)\|^2, \mathbf{z} \sim \mathcal{N}(0,1)$	$\hat{x} = f(\mathbf{z}), \mathbf{z} \sim \mathcal{N}(0,1)$
AR [15] (PMLR'16)	$\mathcal{L} = \|p(x\|\mathbf{z}) - p(x)\|^2, p(x) = \prod_{i=1}^{w \times h} p(x_i\|x_1, x_2, \cdots, x_{i-1}), p(x_1) = \mathcal{N}(0,1)$	$x_i = \prod_{i=1}^{w \times h} p(x_i\|x_1, x_2, \cdots, x_{i-1}), x_1 \sim \mathcal{N}(0,1)$
DM [5] (NIPS'20)	$\mathcal{L}_t = \|f(\sqrt{\overline{\alpha_t}}x_0 + \sqrt{1-\overline{\alpha_t}}\epsilon') - \epsilon'\|^2, \epsilon \sim \mathcal{N}(0,1)$	$x_t \sim \mathcal{N}(\sqrt{1-\beta_t}x_{t-1}, \beta_t \mathbf{I}), \alpha_t = \prod_{s=1}^{t} 1 - \beta_s$

Table 1: This table follows the naming convention in the original papers. p/q originally stands for a distribution. For simplicity, we also interpret p/q as a function in this table. For example, $\mathcal{L}_{rec} = \|p(x\|\mathbf{z}) - p(x)\|^2$ is equivalent to $\mathcal{L}_{rec} = \mathbb{E}_{x \sim p(x), \mathbf{z} \sim q(\mathbf{z}\|x), x' \sim p(x'\|\mathbf{z})} \|x' - x\|^2$. Both of p and q along with G, D, f can be modeled by a neural network.

3.1.3 The Close Form of Forward Process. In the backward process, μ and Σ should be inferred via neural networks. However, in the forward process, careful readers should have noticed that, $q(x_t|x_0)$ is actually a combination of a sequence of known Gaussian distributions $\mathcal{N}(\sqrt{1-\beta_t}x_{t-1}, \beta_t \mathbf{I})$. Therefore, $q(x_t|x_0)$ also obeys a known Gaussian distribution:

$$\alpha_t = 1 - \beta_t;$$
$$\overline{\alpha_t} = \prod_{s=1}^{t} \alpha_s; \quad (3)$$
$$q(x_t|x_0) = \mathcal{N}(\sqrt{\overline{\alpha_t}}x_0, (1 - \overline{\alpha_t})\mathbf{I}).$$

3.1.4 Training. The training loss is simple:

$$\mathcal{L} = \|E(\sqrt{\overline{\alpha_t}}x_0 + \sqrt{1-\overline{\alpha_t}}\epsilon, t) - \epsilon\|^2. \quad (4)$$

E is a noise estimator, usually a U-Net conditioned by t. E aims to predict Gaussian noises ϵ.

3.1.5 Some Details. Σ is simply assumed as scaled $\mathbf{I}$ in pratical, while μ is modeled as follow:

$$\mu(x_t, t) = \frac{1}{\sqrt{\alpha_t}}(x_t - \frac{1-\alpha_t}{\sqrt{1-\overline{\alpha_t}}}E(x_t, t)). \quad (5)$$

In the training stage, a step $t \in [1, T]$ is randomly sampled to otain the loss $\mathcal{L}$.

The generation process must be conducted step by step. The sugguested diffusion steps T should be at least 50 to obtain an okay 128×128 image. For the best image quality, T is usually set to 1000. It indicates that E should be run 1000 times for a generation, consuming about 100 seconds using V100.

3.2 GD

Guided Diffusion (GD) [1] mainly introduces three improvements: **learnable Σ**, **synthesis condition** and **architectural improvement**.

3.2.1 learnable Σ. Instead of fixing Σ at $\mathbf{I}$, GD defines a learnable Σ:

$$\tilde{\beta}_t = \frac{1-\alpha_{t-1}}{1-\alpha_t}\beta_t,$$
$$\Sigma(x_t, t) = \exp(v \log \beta_t + (1-v) \log \tilde{\beta}_t), \quad (6)$$

where v is a learnable parameter. This improvement refines the image and reduces inference time.

3.2.2 Learnable Σ. To obtain images that meet the need of users, we want to add a condition in the synthesis process. For example, the condition is image category constraint. Instead of naively adding a condition input to the U-Net, GD appends a condition head (*i.e.*, a classification head) to the U-Net output. GD train the classification head by classifying the synthetized images, and does NOT need to fine-tune the U-Net parameters. In the test stage, the image synthesis procedure is formulated below:

$$g = \Delta_{x_t} \log(p(y|x_t)),$$
$$x_t \sim \mathcal{N}(\mu + \Sigma g, \Sigma), \quad (7)$$

where $p(y|x_t)$ is a classification CNN. Besides, it should be noticed that U-Net consumes all trainable parameters in DDPM. Therefore, for different conditions, only different condition heads are required.

We use an example to intuitively explain Eq.7. Let our target be a panda image. The original DDPM is able to provide a panda image. However, the chance is poor because of the randomness in Gaussian sampling. Therefore, we "move" the μ to a certain direction that results in a higher propability to provide a panda image. And the gradient g tells what the direction is.

3.2.3 Architectural improvement. The architectural improvements includes:

(1) Increase depth, decreases width.
(2) Increase attention heads.
(3) Single-resolution attention $\rightarrow$ multi-resolution attention.
(4) For upsampling and downsampling, interpolation $\rightarrow$ residual blocks.
(5) Rescaling residual connections with $\frac{1}{\sqrt{2}}$ (not clear yet).

3.3 SGD

Semantic Guidence Diffusion [9] (SGD) is an extension of GD. SGD studies several specific guidances: **language guidance**, **image guidance** and **multimodal guidance**. SGD first generalize the GD condition function of $p(y|x_t)$ to $F(x_t, y, t)$. For different guidances, F is instantiated to a special form.

3.3.1 Language guidance: $F(x_t, l, t) = F''(x_t, t) \cdot F'(l)$. E' encode x_t into an image feature, while $F'(l)$ encode text l into a text feature. The similarity between $F''(x_t, t)$ and $F'(l)$ is measured via dot product.

3.3.2 Image guidance: there are three different image guidances.
The first is semantic-level guidance:

$$F(x_t, x_t', t) = F'(x_t, t) \cdot F'(x_t', t), \quad (8)$$

workflow step	training workflow	test workflow
0	input: caption y, image x	input: caption y
1	CLIP produces CLIP text embedding z_t, CLIP image embedding z_i	CLIP produces CLIP text embedding z_t
2	train a diffusion prior model (DPM) that generates a prior $P(z_i\|y)$ and minimize $\mathbb{E}_{z_i' \sim P(z_i'\|y)}\|\|z_i' - z_i\|\|^2$	DPM generates z_i
3	train a diffusion decoder model (DDM) that inverts z_i into x and minimize $\mathbb{E}_{x' \sim P(x'\|z_i, y)}\|\|x' - x\|\|^2$	DDM inverts z_i into x

Table 2: A simplified DALLE2 training and test workflow.

where x_t' is the reference image noised by Eq. 3. It works similarly with language guidance.

The second is structure-level guidance:

$$F(x_t, x_t', t) = -\sum_j \frac{1}{C_j H_j W_j} \|\|F'(x_t, t)_j - F'(x_t', t)_j\|\|, \qquad (9)$$

where $F'()_j \in \mathcal{R}^{C_j \times H_j \times W_j}$.

The third is style-level guidance:

$$F(x_t, x_t', t) = -\sum_j \|\|Gram(x_t, t) - Gram(x_t', t)\|\|_F^2, \qquad (10)$$

where $Gram$ yields a Gram matrix and $\|\| \cdot \|\|_F$ is Frobenius norm. For the Gram matrix, please refer to https://blog.csdn.net/wangyang20170901/article/details/79037867. For the Frobenius norm, please refer to https://blog.csdn.net/qq_27946691/article/details/56488063.

3.3.3　Multimodel guidance: $F = \sum_{i=1}^g w_i F_i$, where F_i is the $i - th$ model guidance and w_i is a weight. In other words, the multimodel guidance is simply a sum of guidances from different models.

3.4　CFDG

Classifier-Free Diffusion Guidence [6] (CFDG) argues that **1)** GD/SGD is NOT able to jointly train the U-Net and the condition head; **2)** the condition head largely increases inference computation overhead because of the requirement of gradient Δ. Therefore, CFDG integrates the condition into the noise estimator E in Sec. 3.1 and obtains E':

$$E'(x_t, t, y) = E(x_t, t, null) + scale(E(x_t, t, y) - E(x_t, t, null)). \quad (11)$$

However, the whole noise estimator must be re-trained for different condition types.

3.5　GLIDE

GLIDE [10] replaces the category condition of CFDG with language condition (*i.e.*, image caption):

$$E'(x_t, l, l) = E(x_t, t, null) + scale(E(x_t, t, l) - E(x_t, t, null)). \quad (12)$$

Language embedding l is derived from a CLIP model[11].

3.6　DALLE

Although DALLE[13] does not adopt any diffusion model, it is the beginning DALLE generation. The paper is filled with engineering details and is good for future reference. Briefly, there are two stages in DALLE:

(1) DALLE encodes the image into image tokens via a discrete VAE (dVAE);
(2) An AR model is adopted to learn the joint distribution of image and text tokens.

More specifically, DALLE tries to maximize a evidence lower bound:

$$ELB = \mathbb{E}_{z \sim q_\phi(z|x)}(\ln p_\theta(x|y, z) - KL(q_\phi(y, z|x), p_\varphi(y, z))). \quad (13)$$

Here, ϕ is the dVAE encoder; θ is the dVAE decoder; φ is the AR model. x is an images. y is the caption y. z is the image token.

3.7　DALLE2

3.7.1　Two main modules of DALLE2: CLIP image prior and DM.. DALLE2[12] incorerates strengths of CLIP and GLIDE, also consisting of two main modules:

(1) A CLIP image prior $P(z_i|y)$ that produces CLIP[11] image embeddings z_i conditioned on captions y;
(2) A diffusion decoder $P(x|z_i, y)$ that produces images x conditioned on CLIP image embeddings z_i (and optionally captions y).

The DALLE2 workflows is shown in Tab. 2. Note that we will skip CLIP details [11] and only discuss the diffusion decoder..

3.7.2　The diffusion prior model (DPM).. DPM learns $P(z_i|y)$. Ramesh *et. al.* also tries AR to learn $P(z_i|y)$ but find out that the diffusion model is better. The noise estimator of the DPM is a transformer decoder with a causal attention mask. The input sequence consists of: **1)** encoded text l, **2)** the CLIP text embedding z_t, **3)** the diffusion timestep τ and **4)** the noised CLIP image embedding z_i^τ. Different from common DMs, the output directly predicts the *unnoised CLIP image embedding*, instead of the noise. We explain the DPM loss in formula:

$$\mathcal{L}_{prior} = \mathbb{E}_{\tau \sim [1,T], z_i^\tau \sim q_\tau} \|\|f(z_i^\tau, \tau, l, z_t) - z_i\|\|^2. \qquad (14)$$

f is the transformer decoder, q_t is a Gaussian distribution of time step τ. The specific form of q_t is not given in the original paper, and we conjure that it should be similar to Eq. 3.

3.7.3　The diffusion decoder model (DDM).. DDM learns $P(x|z_i, y)$ and is similar to GLIDE. The DALLE2 architecture differs from GLIDE (Eq. 12) mainly in two apects:

(1) The CLIP image embedding z_i is projected and added to a timestep embedding.
(2) The CLIP image embedding z_i is projected into four extra tokens. These tokens are concatenated to output sequence from GLIDE text encoder.

4　CONCLUSION

4.1　Our contributions.

In this paper, we have discussed key ideas of several advanced diffusion models[1, 5, 6, 9, 12] and several other related generative works[2, 3, 7, 10, 13, 15]. Diffusion models take the task of image

Model	test cost	training cost for N condition types	text2image	dataset
DDPM [5]	noise estimation	noise estimation (condition is not allowed)	×	CIFAR10, LSUN
GD [1]	noise estimation + condition gradient derivation	noise estimation + condition derivation $\times N$	×	ImageNet
SGD [9]	noise estimation + condition gradient derivation	noise estimation + condition derivation $\times N$	√	FFHQ, LSUN
CFDG [6]	noise estimation $\times 2$	noise estimation $\times N$	√	ImageNet
GLIDE [6]	noise estimation $\times 2$ + CLIP inference	noise estimation $\times N$	√	a big mixture
DALLE2 [12]	noise estimation $\times (2 + 2)$ + CLIP inference	noise estimation $\times N$	√	a big mixture

Table 3: Resource comparisons. By marking "text2image" as ×, we indicate that the model is not originally proposed for the text-to-image task, but they might be transferred to this task after a proper modification.

Model	repository
DDPM [5] (NIPS'20)	Tensorflow Code https://github.com/hojonathanho/diffusion PyTorch Unofficial Code https://github.com/pesser/pytorch_diffusion
FastDPM [8] (arxiv'21)	Demo https://fastdpm.github.io. PyTorch Code https://github.com/FengNiMa/FastDPM_pytorch
VQ-Diffusion [4] (arxiv'21)	PyTorch Code https://github.com/microsoft/VQ-Diffusion
CFDG [6] (NIPS'21)	Unavailable
CLIP [11] (PMLR'21)	PyTorch Code https://github.com/OpenAI/CLIP
DALLE [13] (PMLR'21)	PyTorch Code for discrete VAE model only. https://github.com/openai/DALLE
v-diffusion sampling	PyTorch Code https://github.com/crowsonkb/v-diffusion-pytorch
DALLE2 [12] (arxiv'22)	A waitlist to reach the demo. https://labs.openai.com/waitlist
DDIM [14] (ICLR'21)	PyTorch Code https://github.com/ermongroup/ddim
GD [1] (NIPS'21)	PyTorch Code https://github.com/openai/guided-diffusion
GLIDE [10] (arxiv'21)	PyTorch Code https://github.com/openai/glide-text2im
SGD [9] (arxiv'21)	Project Page https://xh-liu.github.io/sdg/. Code is unavailable yet.

Table 4: Code repository overview. The code supports training and testing by default. Besides, all of the aboved papers can be reached in Arxiv.

synthesis as a denoising process. In text2image diffusion works[6, 9, 12], the image caption is taken as various conditions to guide the denoising process. Among these works, DALLE2[12] is able to incorperates the strengths of both CLIP and GLIDE to generate high-quality images. We also compare required resources in Tab. 3 and summarize code repositories in Tab. 4 for further study of readers.

4.2 Weaknesses.

Due to time limitation, we failed to include several other diffusion models into our paper[4, 8, 14]. This will be done in the future.

4.3 Future research directions.

Text2image works is now able to produce high-quality images, in which **1)** the objects have clear outlines and reasonable appearances; **2)** object relations can match the text description. Therefore, their might be an approach to produce a object mask that can tell different objects, or tell the foreground object so that the model can be adopted to more comprehensively researches or applications.

Existing works are not sensitive to numbers. For example, a text of "two cars are running on the street" might not be able to result in an image of two cars. On the contrary, humans can well distinguish two cars from one or three. To enhance numerical sensitivity might be a good research direction.

Step-by-step sampling is time-consuming. Efficiency is thus not an advantage of diffusion models. To reduce inference time will also be a promising research direction. It could be accomplished by reducing sampling steps or a coarse-to-fine generation framework,

REFERENCES

[1] Prafulla Dhariwal and Alexander Nichol. 2021. Diffusion models beat gans on image synthesis. *Advances in Neural Information Processing Systems* 34 (2021).
[2] Laurent Dinh, David Krueger, and Yoshua Bengio. 2014. Nice: Non-linear independent components estimation. *ICLR workshop* (2014).
[3] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. 2014. Generative adversarial nets. *Advances in neural information processing systems* 27 (2014).
[4] Shuyang Gu, Dong Chen, Jianmin Bao, Fang Wen, Bo Zhang, Dongdong Chen, Lu Yuan, and Baining Guo. 2021. Vector quantized diffusion model for text-to-image synthesis. *arXiv preprint arXiv:2111.14822* (2021).
[5] Jonathan Ho, Ajay Jain, and Pieter Abbeel. 2020. Denoising diffusion probabilistic models. *Advances in Neural Information Processing Systems* 33 (2020), 6840–6851.
[6] Jonathan Ho and Tim Salimans. 2021. Classifier-Free Diffusion Guidance. In *NeurIPS 2021 Workshop on Deep Generative Models and Downstream Applications*.
[7] Diederik P Kingma and Max Welling. 2014. Auto-encoding variational bayes. *ICLR* (2014).
[8] Zhifeng Kong and Wei Ping. 2021. On fast sampling of diffusion probabilistic models. *arXiv preprint arXiv:2106.00132* (2021).
[9] Xihui Liu, Dong Huk Park, Samaneh Azadi, Gong Zhang, Arman Chopikyan, Yuxiao Hu, Humphrey Shi, Anna Rohrbach, and Trevor Darrell. 2021. More

Control for Free! Image Synthesis with Semantic Diffusion Guidance. *arXiv preprint arXiv:2112.05744* (2021).

[10] Alex Nichol, Prafulla Dhariwal, Aditya Ramesh, Pranav Shyam, Pamela Mishkin, Bob McGrew, Ilya Sutskever, and Mark Chen. 2021. Glide: Towards photorealistic image generation and editing with text-guided diffusion models. *arXiv preprint arXiv:2112.10741* (2021).

[11] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. 2021. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning*. PMLR, 8748–8763.

[12] Aditya Ramesh, Prafulla Dhariwal, Alex Nichol, Casey Chu, and Mark Chen. 2022. Hierarchical Text-Conditional Image Generation with CLIP Latents. *arXiv preprint arXiv:2204.06125* (2022).

[13] Aditya Ramesh, Mikhail Pavlov, Gabriel Goh, Scott Gray, Chelsea Voss, Alec Radford, Mark Chen, and Ilya Sutskever. 2021. Zero-shot text-to-image generation. In *International Conference on Machine Learning*. PMLR, 8821–8831.

[14] Jiaming Song, Chenlin Meng, and Stefano Ermon. 2021. Denoising diffusion implicit models. (2021).

[15] Aaron Van Oord, Nal Kalchbrenner, and Koray Kavukcuoglu. 2016. Pixel recurrent neural networks. In *International conference on machine learning*. PMLR, 1747–1756.